

LINUX KERNEL ARCHITECTURE

A DEEP DIVE INTO THE CORE OF THE OS

1. USER SPACE (APPLICATIONS & GLIBC):

- CONTAINS USER APPS AND THE GNU C LIBRARY (GLIBC).
- GLIBC ACTS AS THE PRIMARY WRAPPER FOR SYSTEM CALLS.

2. SYSTEM CALL INTERFACE (SCI):

- THE GATEKEEPER. IT TRANSITIONS REQUESTS FROM USER SPACE TO KERNEL SPACE.
- ENSURES SECURITY AND PROVIDES A STABLE API FOR DEVELOPERS.

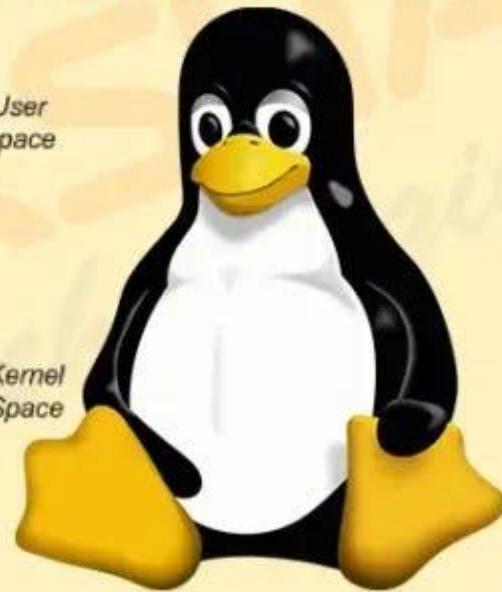
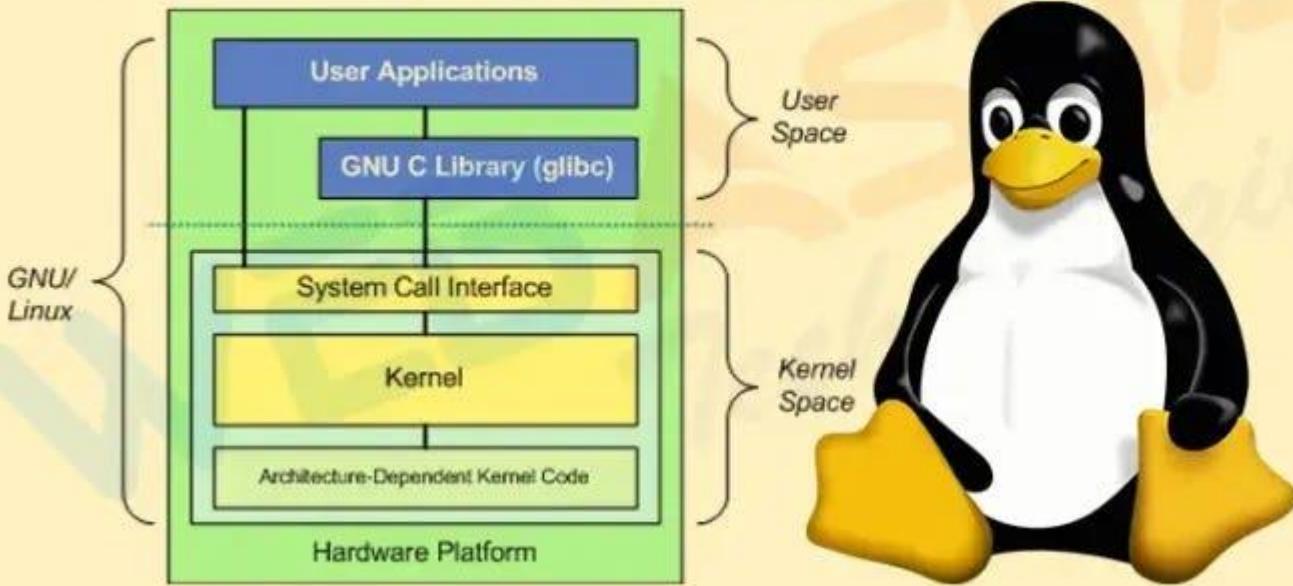
3. KERNEL SPACE (THE ENGINE):

- THE 'MAIN KERNEL' MANAGES SCHEDULING, MEMORY, AND FILESYSTEMS.
- 'ARCH-DEPENDENT CODE' OPTIMIZES THE KERNEL FOR SPECIFIC CPUS (x86, ARM).

4. HARDWARE PLATFORM:

- THE PHYSICAL LAYER (CPU, RAM, DISK, NIC).
- CONTROLLED BY THE KERNEL VIA DEVICE DRIVERS.

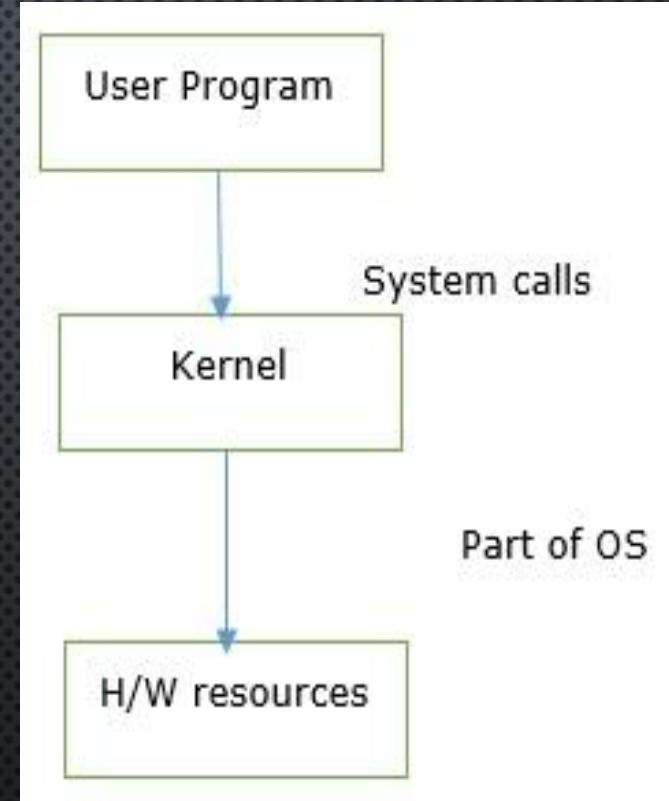
What is Linux & Linux Kernel?



Architectural Breakdown: User vs. Kernel Space

SYSTEM ARCHITECTURE DETAILS:

- TRAP MECHANISM: When glibc executes a System Call, it triggers a 'Software Interrupt' (Trap), forcing the CPU to switch from Ring 3 (User) to Ring 0 (Kernel) mode.
- MONOLITHIC DESIGN: Unlike Microkernels, Linux runs the Scheduler, VFS, and Network Stack within a single address space to eliminate IPC overhead and boost speed.
- ARCH-DEPENDENT LAYER: This is the 'Hardware Abstraction Layer' (HAL). It contains assembly code specific to CPUs (like x86_64 or ARM), allowing the core kernel to remain portable.
- GLIBC WRAPPER: Applications don't talk to the kernel; they talk to glibc. This ensures that if the kernel API changes, the application code doesn't have to.



PROCESS SCHEDULER

- MANAGES CPU TIME FOR ALL PROCESSES.
- USES THE COMPLETELY FAIR SCHEDULER (CFS).
- SUPPORTS PREEMPTION AND MULTI-CORE SCALABILITY.

VIRTUAL FILE SYSTEM (VFS)

- AN ABSTRACTION LAYER FOR STORAGE.
- SUPPORTS MULTIPLE FORMATS (EXT4, XFS, BTRFS).
- STANDARDIZES 'READ/WRITE' OPERATIONS FOR APPLICATIONS.

NETWORK STACK

- IMPLEMENTS TCP/IP AND OTHER PROTOCOLS.
- USES A SOCKET-BASED API FOR COMMUNICATION.
- INTERFACES DIRECTLY WITH NETWORK INTERFACE CARDS (NICs).