

# Implement a Custom System Call using Docker

## Step -1:

Install docker on the host machine/vm

```
sudo apt update
```

```
sudo apt install -y docker.io qemu-system-x86
```

verify whether docker has been installed in the terminal

```
docker --version
```

## Step -2:

Create docker image for kernel build

i)create a project directory

(e.g) kernel-docker.

```
student@student-virtual-machine:~/kernel-docker$ mkdir kernel-docker
student@student-virtual-machine:~/kernel-docker$ cd kernel-docker
```

## Step -3:

Creating a Dockerfile and building it.

```
student@student-virtual-machine:~/kernel-docker$ vi Dockerfile
student@student-virtual-machine:~/kernel-docker$ docker build -t kernel-builder

FROM ubuntu:22.04
ENV DEBIAN_FRONTEND=noninteractive
RUN apt update && apt install -y \
    build-essential \
    bc \
    bison \
    flex \
    libssl-dev \
    libelf-dev \
    libncurses-dev \
    git \
    vim \
    wget \
    cpio \
    && rm -rf /var/lib/apt/lists/*
WORKDIR /kernel
```

## Step -4:

To run the kernel builder container

```
student@student-virtual-machine:~/kernel-docker$ docker run -it --name kernel-dev \
-v $PWD:/kernel \
kernel-builder
```

## Step -5:

Download linux kernel source code in the docker and extract it using tar cmd

```
root@92f206db8b9a:/kernel# wget https://cdn.kernel.org/pub/linux/kernel/v6.x/linux-6.6.tar.xz
--2026-02-03 16:05:01-- https://cdn.kernel.org/pub/linux/kernel/v6.x/linux-6.6.tar.xz
```

After installing navigate to linux-6.6 directory

```
root@92f206db8b9a:/kernel# cd linux-6.6
```


## Step -6:

Adding custom system call

i)create a folder in kernel directory of linux-6.6

```
root@92f206db8b9a:/kernel/linux-6.6# mkdir kernel/custom_syscalls
root@92f206db8b9a:/kernel/linux-6.6# vim kernel/custom_syscalls/stringrev_syscall.c
```

ii)create a program file for implementing custom system call for string reverse



```
1 #include <linux/kernel.h>
2 #include <linux/syscalls.h>
3 #include <linux/uaccess.h>
4
5 SYSCALL_DEFINE3(reverse_string,
6     char __user *, input,
7     char __user *, output,
8     int, len)
9 {
10     char kbuf[256];
11     int i;
12
13     if (len <= 0 || len > 255)
14         return -EINVAL;
15     if (copy_from_user(kbuf, input, len))
16         return -EFAULT;
17     kbuf[len] = '\0';
18     for (i = 0; i < len / 2; i++) {
19         char temp = kbuf[i];
20         kbuf[i] = kbuf[len - i - 1];
21         kbuf[len - i - 1] = temp;
22     }
23     if (copy_to_user(output, kbuf, len + 1))
24         return -EFAULT;
25     printk(KERN_INFO "reverse_string syscall executed\n");
26     return 0;
27 }
```

## Step -7:

create a makefile for custom system call program

```
root@92f206db8b9a:/kernel/linux-6.6# vim kernel/custom_syscalls/Makefile
```

In makefile add this

```
obj-y := stringrev_syscall.o
```

Then Add a entry in makefile for the custom system call in kernel makefile

```
root@92f206db8b9a:/kernel/linux-6.6# vim kernel/Makefile
```

```
obj-y += custom_syscalls/
```

## Step -8:

Register the syscall number in (x86\_64) tbl file

```
root@92f206db8b9a:/kernel/linux-6.6# vim arch/x86/entry/syscalls/syscall_64.tbl
```

Add this line

```
454 common reverse_string sys_reverse_string
```

## Step -9:

Declare syscall prototype in syscalls.h header file.

```
root@92f206db8b9a:/kernel/linux-6.6# vim include/linux/syscalls.h
```

```
asmlinkage long sys_reverse_string(const char __user *input, char __user *output, int len);
```

## Step -10:

Configure the kernel now

```
root@92f206db8b9a:/kernel/linux-6.6# make defconfig
```

## Step -11:

Build the kernel inside the docker

```
root@92f206db8b9a:/kernel/linux-6.6# make -j$(nproc)
```

```
Kernel: arch/x86/boot/bzImage is ready (#1)
```

## Step -12:

Now create a minimal root file system

i)create rootfs folder in kernel-docker directory and navigate to it and create necessary directories inside it

```
mkdir rootfs  
cd rootfs  
mkdir -p bin sbin etc proc sys usr/bin usr/sbin
```

ii)create init process

1)create a init file

```
vim init
```

2)add this line inside

```
#!/bin/sh
mount -t proc none /proc
mount -t sysfs none /sys
exec /bin/sh
```

3) make the file executable

```
chmod +x init
```

### Step -13:

Install busybox in host or docker

```
sudo apt install -y busybox-static
```

Copy busybox to rootfs directory

```
cp /bin/busybox rootfs/bin/
```

Now create applet links for required one with busybox in rootfs

```
ln -s busybox rootfs/bin/sh
```

```
ln -s busybox rootfs/bin/cd
```

```
ln -s busybox rootfs/bin/dmesg
```

```
ln -s busybox rootfs/bin/head
```

```
ln -s busybox rootfs/bin/tail
```

```
ln -s busybox rootfs/bin/mkdir
```

now build initramfs from rootfs directory

```
cd rootfs
find . | cpio -o --format=newc > ../initramfs.cpio
```

## Step -14:

Test the custom system call using a program

```
#include <stdio.h>
#include <unistd.h>
#include <string.h>
#include <sys/syscall.h>

#define __NR_reverse_string 454

int main()
{
    char input[256];
    char output[256];
    int len;

    scanf("%s", input);
    len = strlen(input);

    long ret = syscall(__NR_reverse_string, input, output, len);

    if(ret < 0){
        perror("syscall failed");
        return 1;
    }

    if(ret >= 0)
        write(1, "Reversed: ", 10);
        write(1, output, strlen(output));
        write(1, "\n", 1);

    return 0;
}
```

Compile the program in static mode

```
gcc test.c -static -o test
```

move the obj file to kernel-docker/rootfs and create subdirectory there and store it.

```
mv ~/kernel-docker/test ~/kernel-docker/rootfs/test
```

now once again build initramfs from rootfs directory

```
cd rootfs
find . | cpio -o --format=newc > ../initramfs.cpio
```

## Step -15:

Boot kernel using QEMU (Host)

Navigate to your folder you created in rootfs for obj file and run it to test the custom system call.

```
BusyBox v1.30.1 (Ubuntu 1:1.30.1-7ubuntu3.1) built-in shell (ash)
Enter 'help' for a list of built-in commands.

/bin/sh: can't access tty; job control turned off
/ # cd fold
/fold # ./test
Anirudh
[ 42.396118] reverse_string syscall executed
Reversed: hdurinA
```

Then type dmesg cmd to check the kernel buffer output during custom system call

```
[ 42.396118] reverse_string syscall executed
[ 42.400652] test (50) used greatest stack depth: 14544 bytes left
[ 51.762417] dmesg (51) used greatest stack depth: 14096 bytes left
```